


# Лекция 3

## СОДЕРЖАНИЕ

- *разделяемая память (shared memory)*


## L1 кэш память



## Разделяемая память (shared memory)

Разделяемая память CUDA – память с низкой латентностью и высокой пропускной способностью.

Высокая пропускная способность обеспечивается параллельным выполнением запросов, благодаря разделению памяти на отдельные модули, банки памяти.



**Если более одной нити варпа обращаются к одному и тому же банку, то происходит конфликт, который разрешается сериализацией выполнения запроса.**

# Выделение разделяемой памяти

Разделяемая память выделяется (статически или динамически) только на устройстве. Область видимости – нити одного блока. Время жизни – время выполнения ядра.

## Статическое выделение:

```
#define N 3
#define M 512
__global__ void gTest(){
    __shared__ float s[N][M];
    .....
}
```

## Динамическое выделение:

```
extern __shared__ float s[];
__global__ void gTest2(){
    float* a=(float*)s;
    float* b=(float*)&s[512];
    float* c=(float*)&s[1024];
    .....
}
```

```
gTest2<<<100,32,N*M>>>();
```



3-й параметр – размер разделяемой памяти.

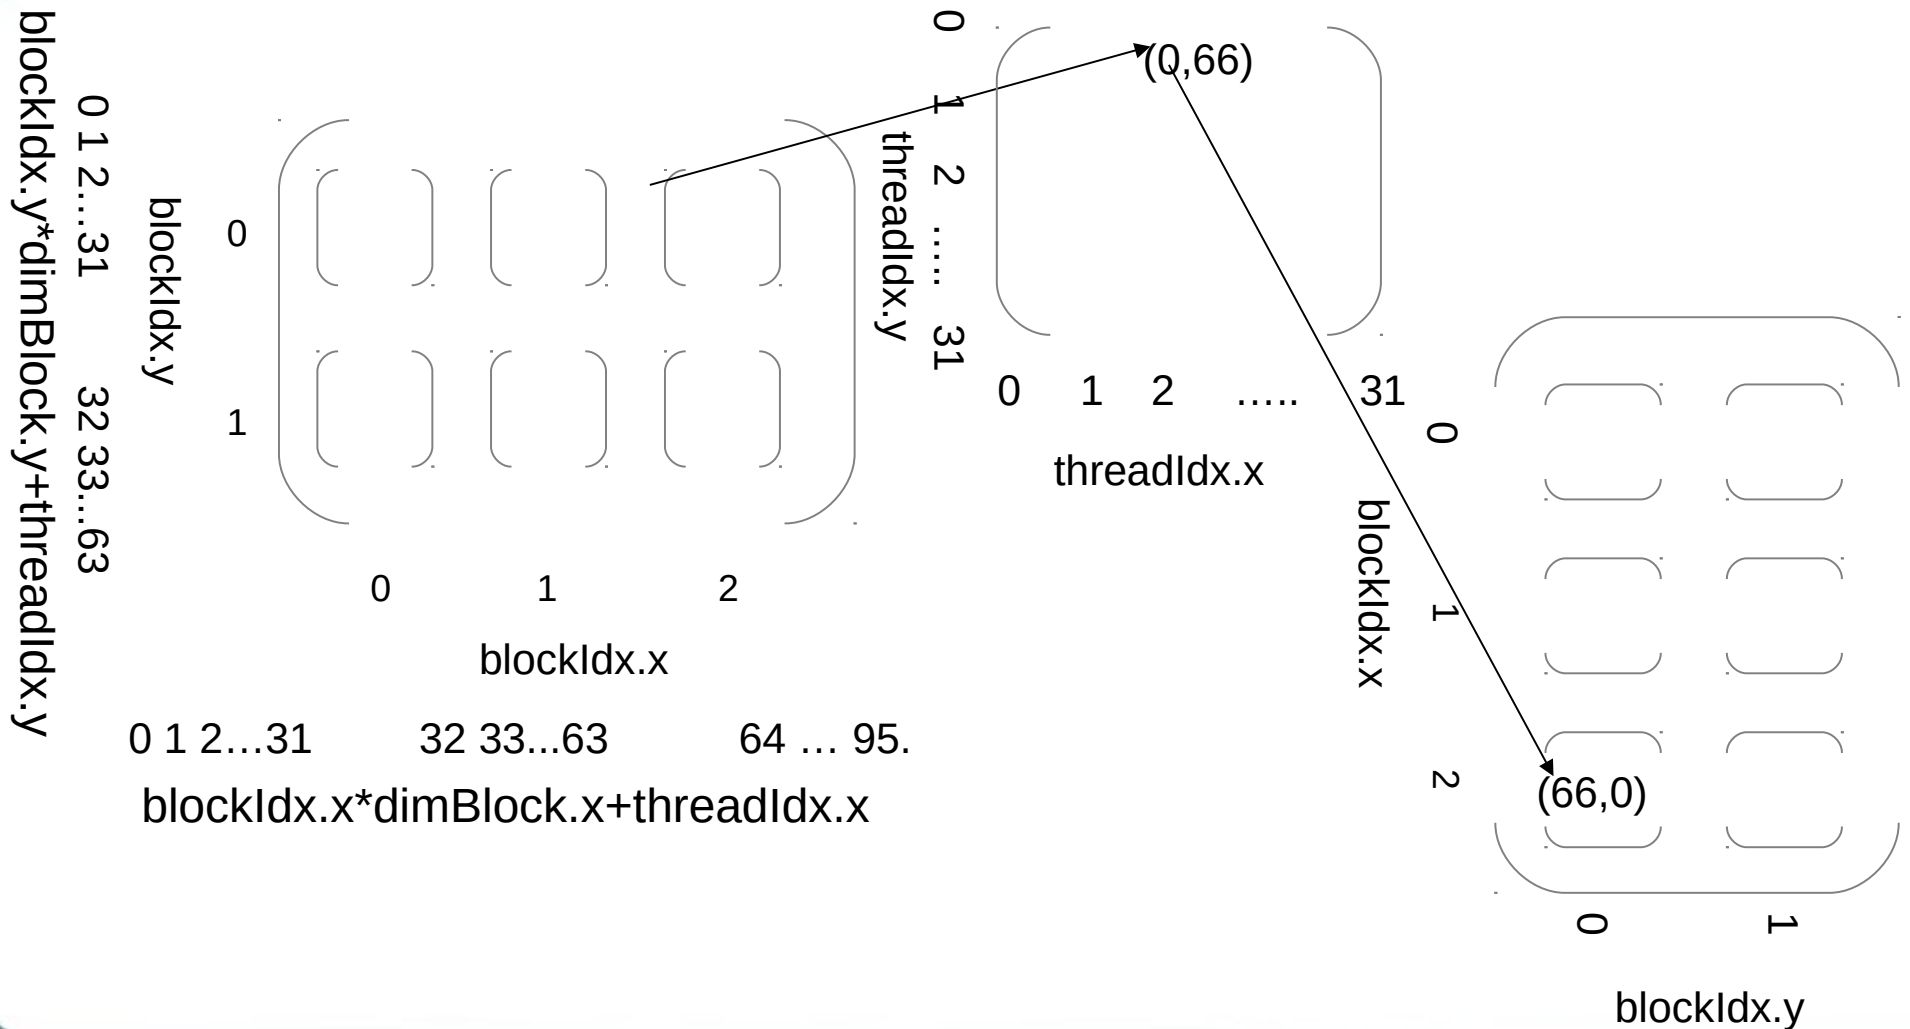
## Задание (лабораторная 4)

Оптимизировать операцию транспонирования матриц используя *coalescing* и *shared memory*.

## Инициализация матрицы

```
__global__ void glInitializeStorage(float* storage_d){  
    int i=threadIdx.x+blockIdx.x*blockDim.x;  
    int j=threadIdx.y+blockIdx.y*blockDim.y;  
    int N=blockDim.x*gridDim.x;  
  
    storage_d[i+j*N]=i+j*N;  
}
```

# Простое транспонирование (схема)




## Простое транспонирование

```
__global__ void gTranspose0(float* storage_d, float* storage_d_t){  
    int i=threadIdx.x+blockIdx.x*blockDim.x;  
    int j=threadIdx.y+blockIdx.y*blockDim.y;  
    int N=blockDim.x*gridDim.x;  
  
    storage_d_t[j+i*N]=storage_d[i+j*N];  
}
```



# Использование shared memory с конфликтом банков (схема)



## Использование shared memory с конфликтом банков

```
#define SH_DIM 32

__global__ void gTranspose1(float* storage_d, float* storage_d_t){
    __shared__ float buffer[SH_DIM][SH_DIM];

    int i=threadIdx.x+blockIdx.x*blockDim.x;
    int j=threadIdx.y+blockIdx.y*blockDim.y;
    int N=blockDim.x*gridDim.x;

    buffer[threadIdx.y][threadIdx.x]=storage_d[i+j*N];
    __syncthreads();

    i=threadIdx.x+blockIdx.y*blockDim.x;
    j=threadIdx.y+blockIdx.x*blockDim.y;

    storage_d_t[i+j*N]=buffer[threadIdx.x][threadIdx.y];
}
```


## Использование shared memory с разрешением конфликтов банков

```
__global__ void gTranspose2(float* storage_d, float* storage_d_t){  
    __shared__ float buffer[SH_DIM][SH_DIM+1];  
  
    int i=threadIdx.x+blockIdx.x*blockDim.x;  
    int j=threadIdx.y+blockIdx.y*blockDim.y;  
    int N=blockDim.x*gridDim.x;  
  
    buffer[threadIdx.y][threadIdx.x]=storage_d[i+j*N];  
    __syncthreads();  
  
    i=threadIdx.x+blockIdx.y*blockDim.x;  
    j=threadIdx.y+blockIdx.x*blockDim.y;  
  
    storage_d_t[i+j*N]=buffer[threadIdx.x][threadIdx.y];  
}
```

# Как избежать конфликта банков разделяемой памяти

`__shared__ float buffer[SH_DIM][SH_DIM+1];`

Размещение массива `buffer` в разделяемой памяти (shared memory):



# Конфигурация нитей и запуск ядра

```
.....  
int main(int argc, char* argv){  
    if(argc<2){  
        fprintf(stderr, "USAGE: matrix <dimension of matrix>\n");  
        return -1;  
    }  
    int N=atoi(argv[1]);  
  
    const int max_size=1024;  
    int dim_of_threads=SH_DIM;  
    int size=N/dim_of_threads+(N%dim_of_threads>0);  
    int dim_of_blocks=(size>max_size)?max_size:size;  
    .....  
    float *storage_d, *storage_d_t, *storage_h;
```

## Конфигурация нитей и запуск ядра

```
cudaMalloc((void**)&storage_d, N*N*sizeof(float));  
cudaMalloc((void**)&storage_d_t, N*N*sizeof(float));  
storage_h=(float*)calloc(N*N, sizeof(float));  
  
glInitializeStorage<<<dim3(dim_of_blocks,dim_of_blocks),  
                        dim3(dim_of_threads,dim_of_threads)>>>(storage_d,N);  
cudaDeviceSynchronize();  
  
cudaMemcpy(storage_h, storage_d, N*N*sizeof(float),  
                                                    cudaMemcpyDeviceToHost);  
  
Output(storage_h, N);
```

## Конфигурация нитей и запуск ядра

```
gTranspose0<<<dim3(dim_of_blocks, dim_of_blocks),  
               dim3(dim_of_threads, dim_of_threads)>>>  
               (storage_d, storage_d_t, N);  
cudaDeviceSynchronize();  
  
gTranspose1<<<dim3(dim_of_blocks, dim_of_blocks),  
               dim3(dim_of_threads, dim_of_threads)>>>  
               (storage_d, storage_d_t, N);  
CudaDeviceSynchronize();  
  
gTranspose2<<<dim3(dim_of_blocks, dim_of_blocks),  
               dim3(dim_of_threads, dim_of_threads)>>>  
               (storage_d, storage_d_t, N);  
cudaDeviceSynchronize();
```

## Конфигурация нитей и запуск ядра

```
cudaMemcpy(storage_h, storage_d_t, N*N*sizeof(float),  
                                                    cudaMemcpyDeviceToHost);  
Output(storage_h, N);  
  
cudaFree(storage_d);  
cudaFree(storage_d_t);  
free(storage_h);  
  
return 0;  
}
```



# Время выполнения ядра с разными архитектурами CUDA (GeForce 560Ti, *Fermi*, CUDA 7.5, *compute capabilities 2.1*)

```
malkov@linux-5002:~/WORKSHOP/EDUCATION/SibSUTIS/COURSES/2016-2017/CUDA/Lectures/Lecture4/Lab4-2> nvcc -arch=sm_21 tran_matr.cu -o tran_matr
malkov@linux-5002:~/WORKSHOP/EDUCATION/SibSUTIS/COURSES/2016-2017/CUDA/Lectures/Lecture4/Lab4-2> nvprof ./tran_matr 8192
==6208== NVPROF is profiling process 6208, command: ./tran_matr 8192
    0          2048          4096          6144
16777216      16779264      16781312      16783360
33554432      33556480      33558528      33560576
50331648      50333696      50335744      50337792

    0          16777216      33554432      50331648
2048          16779264      33556480      50333696
4096          16781312      33558528      50335744
6144          16783360      33560576      50337792

==6208== Profiling application: ./tran_matr 8192
==6208== Profiling result:
Time(%)      Time      Calls      Avg      Min      Max      Name
49.45%    111.14ms        2    55.572ms  44.880ms  66.263ms  [CUDA memcpy DtoH]
17.18%    38.623ms        1    38.623ms  38.623ms  38.623ms  gTranspose1(float*, float*, int)
15.37%    34.543ms        1    34.543ms  34.543ms  34.543ms  gTranspose0(float*, float*, int)
12.00%    26.967ms        1    26.967ms  26.967ms  26.967ms  gTranspose2(float*, float*, int)
 6.01%    13.503ms        1    13.503ms  13.503ms  13.503ms  gInitializeStorage(float*, int)

==6208== API calls:
Time(%)      Time      Calls      Avg      Min      Max      Name
38.77%    114.78ms        4    28.696ms  13.503ms  38.661ms  cudaThreadSynchronize
37.71%    111.65ms        2    55.827ms  45.073ms  66.580ms  cudaMemcpy
23.25%    68.847ms        2    34.423ms  310.52us  68.536ms  cudaMalloc
 0.13%    375.60us       83    4.5250us    225ns    164.34us  cuDeviceGetAttribute
 0.08%    245.58us        2    122.79us  96.540us  149.04us  cudaFree
 0.03%    81.900us        4    20.475us  7.3030us  27.520us  cudaLaunch
 0.02%    52.788us        1    52.788us  52.788us  52.788us  cuDeviceTotalMem
 0.01%    39.152us        1    39.152us  39.152us  39.152us  cuDeviceGetName
 0.00%    7.3490us       11      668ns    134ns    3.3850us  cudaSetupArgument
 0.00%    4.6030us        4    1.1500us  347ns    1.5880us  cudaConfigureCall
 0.00%    1.7530us        2      876ns    396ns    1.3570us  cuDeviceGetCount
 0.00%      836ns        2      418ns    264ns     572ns  cuDeviceGet
```

## Время выполнения ядра с разными архитектурами CUDA (GeForce 560Ti, *Fermi*, CUDA 7.5, *compute capabilities 2.1*)

```
/Lab4-2> nvcc -arch=sm_21 tran_matr.cu -o tran_matr
/Lab4-2> nvprof ./tran_matr 8192
```

.....

Avg	Min	Max	Name
55.572ms	44.880ms	66.263ms	[CUDA memcpy DtoH]
38.623ms	38.623ms	38.623ms	gTranspose1(float*, float*, int)
34.543ms	34.543ms	34.543ms	gTranspose0(float*, float*, int)
26.967ms	26.967ms	26.967ms	gTranspose2(float*, float*, int)
13.503ms	13.503ms	13.503ms	gInitializeStorage(float*, int)

# Время выполнения ядра с разными архитектурами CUDA (GeForce 650Ti, *Kepler*, CUDA 6.5, *compute capabilities 2.1*)

```
malkov@dew: ~/WORKSHOP/PROJECTS/CUDA-GDB/CUDA_GDB> nvcc -arch=sm_21 tran_matr.cu -o tran_matr
malkov@dew: ~/WORKSHOP/PROJECTS/CUDA-GDB/CUDA_GDB> nvprof ./tran_matr 8192
==4657== NVPROF is profiling process 4657, command: ./tran_matr 8192

      0      2048      4096      6144
16777216      16779264      16781312      16783360
33554432      33556480      33558528      33560576
50331648      50333696      50335744      50337792

      0      16777216      33554432      50331648
2048      16779264      33556480      50333696
4096      16781312      33558528      50335744
6144      16783360      33560576      50337792

==4657== Profiling application: ./tran_matr 8192
==4657== Profiling result:
Time(%)      Time      Calls      Avg      Min      Max      Name
43.93%      98.649ms      2      49.325ms      43.201ms      55.449ms      [CUDA memcpy DtoH]
19.27%      43.280ms      1      43.280ms      43.280ms      43.280ms      gTranspose0(float*, float*, int)
16.90%      37.940ms      1      37.940ms      37.940ms      37.940ms      gTranspose1(float*, float*, int)
13.71%      30.785ms      1      30.785ms      30.785ms      30.785ms      gTranspose2(float*, float*, int)
6.19%      13.890ms      1      13.890ms      13.890ms      13.890ms      gInitializeStorage(float*, int)

==4657== API calls:
Time(%)      Time      Calls      Avg      Min      Max      Name
48.14%      125.90ms      4      31.476ms      13.892ms      43.282ms      cudaThreadSynchronize
37.90%      99.118ms      2      49.559ms      43.352ms      55.766ms      cudaMemcpy
13.69%      35.801ms      2      17.901ms      177.24us      35.624ms      cudaMalloc
0.12%      314.36us      2      157.18us      138.16us      176.19us      cudaFree
0.10%      266.18us      83      3.2070us      143ns      120.28us      cuDeviceGetAttribute
0.03%      75.920us      4      18.980us      7.0430us      30.630us      cudaLaunch
0.01%      20.542us      1      20.542us      20.542us      20.542us      cuDeviceTotalMem
0.01%      16.244us      1      16.244us      16.244us      16.244us      cuDeviceGetName
0.00%      6.3660us      11      578ns      131ns      3.5570us      cudaSetupArgument
0.00%      2.6520us      4      663ns      237ns      1.2360us      cudaConfigureCall
0.00%      1.2360us      2      618ns      245ns      991ns      cuDeviceGetCount
0.00%      544ns      2      272ns      208ns      336ns      cuDeviceGet
```

## Время выполнения ядра с разными архитектурами CUDA (GeForce 650Ti, *Kepler*, CUDA 6.5, *compute capabilities 2.1*)

```
B/CUDA_GDB> nvcc -arch=sm_21 tran_matr.cu -o tran_matr
B/CUDA_GDB> nvprof ./tran_matr 8192
```

```
.....
      Avg      Min      Max  Name
49.325ms  43.201ms  55.449ms  [CUDA memcpy DtoH]
43.280ms  43.280ms  43.280ms  gTranspose0(float*, float*, int)
37.940ms  37.940ms  37.940ms  gTranspose1(float*, float*, int)
30.785ms  30.785ms  30.785ms  gTranspose2(float*, float*, int)
13.890ms  13.890ms  13.890ms  gInitializeStorage(float*, int)
```

# Время выполнения ядра с разными архитектурами CUDA (Fermi, CUDA 6.5, compute capabilities 1.2)

```
malkov@dew:~/WORKSHOP/PROJECTS/CUDA-GDB/CUDA_GDB> nvcc -arch=sm_12 tran_matr.cu -o tran_matr
nvcc warning : The 'compute_11', 'compute_12', 'compute_13', 'sm_11', 'sm_12', and 'sm_13' architectures are deprecated, and may be removed in a future release.
malkov@dew:~/WORKSHOP/PROJECTS/CUDA-GDB/CUDA_GDB> nvprof ./tran_matr 8192
==4725== NVPROF is profiling process 4725, command: ./tran_matr 8192
  0          2048          4096          6144
16777216     16779264     16781312     16783360
33554432     33556480     33558528     33560576
50331648     50333696     50335744     50337792

  0          16777216     33554432     50331648
2048         16779264     33556480     50333696
4096         16781312     33558528     50335744
6144         16783360     33560576     50337792

==4725== Profiling application: ./tran_matr 8192
==4725== Profiling result:
Time(%)      Time      Calls      Avg      Min      Max      Name
40.49%    98.962ms        2    49.481ms  42.947ms  56.015ms  [CUDA memcpy DtoH]
19.03%    46.521ms        1    46.521ms  46.521ms  46.521ms  gTranspose0(float*, float*, int)
18.26%    44.636ms        1    44.636ms  44.636ms  44.636ms  gTranspose1(float*, float*, int)
15.04%    36.757ms        1    36.757ms  36.757ms  36.757ms  gTranspose2(float*, float*, int)
 7.18%    17.547ms        1    17.547ms  17.547ms  17.547ms  gInitializeStorage(float*, int)

==4725== API calls:
Time(%)      Time      Calls      Avg      Min      Max      Name
45.60%    145.47ms        4    36.368ms  17.549ms  46.523ms  cudaThreadSynchronize
31.32%    99.911ms        2    49.956ms  43.101ms  56.810ms  cudaMemcpy
22.85%    72.891ms        2    36.446ms  202.74us  72.688ms  cudaMalloc
 0.10%    321.37us        2    160.69us  142.06us  179.31us  cudaFree
 0.09%    275.91us       83    3.3240us   150ns   126.87us  cuDeviceGetAttribute
 0.02%    79.513us        4    19.878us   7.3050us  32.134us  cudaLaunch
 0.01%    20.699us        1    20.699us  20.699us  20.699us  cuDeviceTotalMem
 0.01%    16.680us        1    16.680us  16.680us  16.680us  cuDeviceGetName
 0.00%    6.9210us       11     629ns   132ns    3.6860us  cudaSetupArgument
 0.00%    2.7000us        4     675ns   234ns    1.1750us  cudaConfigureCall
 0.00%    1.4430us        2     721ns   260ns    1.1830us  cuDeviceGetCount
 0.00%      491ns        2     245ns   216ns    275ns    cuDeviceGet
```

## Время выполнения ядра с разными архитектурами CUDA (*Fermi*, CUDA 6.5, *compute capabilities 1.2*)

```
/CUDA_GDB> nvcc -arch=sm_12 tran_matr.cu -o tran_matr
```

```
.....
```

Avg	Min	Max	Name
49.481ms	42.947ms	56.015ms	[CUDA memcpy DtoH]
46.521ms	46.521ms	46.521ms	gTranspose0(float*, float*, int)
44.636ms	44.636ms	44.636ms	gTranspose1(float*, float*, int)
36.757ms	36.757ms	36.757ms	gTranspose2(float*, float*, int)
17.547ms	17.547ms	17.547ms	gInitializeStorage(float*, int)

# Время выполнения ядра с разными архитектурами CUDA (Tesla K40m, *Kepler*, CUDA 7.5, *compute capabilities 3.5*)

```
malkov@master:~/WORKSPACE/cuda-gdb/ nvcc -arch=sm_35 tran_matr.cu -o tran_matr
malkov@master:~/WORKSPACE/cuda-gdb/ qsub -I
qsub: waiting for job 243.master to start
qsub: job 243.master ready

malkov@n01:~/ cd WORKSPACE/cuda-gdb/
malkov@n01:~/WORKSPACE/cuda-gdb/ nvprof ./tran_matr 8192
==15589== NVPROF is profiling process 15589, command: ./tran_matr 8192
0          2048          4096          6144
16777216   16779264   16781312   16783360
33554432   33556480   33558528   33560576
50331648   50333696   50335744   50337792

0          16777216   33554432   50331648
2048       16779264   33556480   50333696
4096       16781312   33558528   50335744
6144       16783360   33560576   50337792

==15589== Profiling application: ./tran_matr 8192
==15589== Profiling result:
Time(%)    Time      Calls      Avg      Min      Max      Name
85.79%    147.53ms      2    73.764ms  26.145ms  121.38ms  [CUDA memcpy DtoH]
4.86%     8.3598ms      1    8.3598ms  8.3598ms  8.3598ms  gTranspose0(float*, float*, int)
4.37%     7.5069ms      1    7.5069ms  7.5069ms  7.5069ms  gTranspose1(float*, float*, int)
3.48%     5.9834ms      1    5.9834ms  5.9834ms  5.9834ms  gTranspose2(float*, float*, int)
1.50%     2.5776ms      1    2.5776ms  2.5776ms  2.5776ms  gInitializeStorage(float*, int)

==15589== API calls:
Time(%)    Time      Calls      Avg      Min      Max      Name
60.35%    265.25ms      2    132.62ms  443.24us  264.80ms  cudaMalloc
33.77%    148.42ms      2     74.210ms  26.253ms  122.17ms  cudaMemcpy
5.56%     24.436ms      4     6.1090ms  2.5801ms  8.3542ms  cudaThreadSynchronize
0.18%     795.15us    166    4.7900us      243ns  177.00us  cuDeviceGetAttribute
0.07%     310.88us      2    155.44us  130.99us  179.88us  cudaFree
0.02%     102.07us      4     25.516us  7.7810us  52.205us  cudaLaunch
0.02%     101.68us      2     50.837us  37.004us  64.671us  cuDeviceGetName
0.02%     100.57us      2     50.284us  47.734us  52.834us  cuDeviceTotalMem
0.00%      9.3680us     11      851ns      141ns  5.6560us  cudaSetupArgument
0.00%      4.8890us      2     2.4440us  1.6440us  3.2450us  cuDeviceGetCount
0.00%      3.8460us      4      961ns      244ns  2.3180us  cudaConfigureCall
0.00%      2.0930us      4      523ns      441ns    645ns  cuDeviceGet

malkov@n01:~/WORKSPACE/cuda-gdb/ █
```



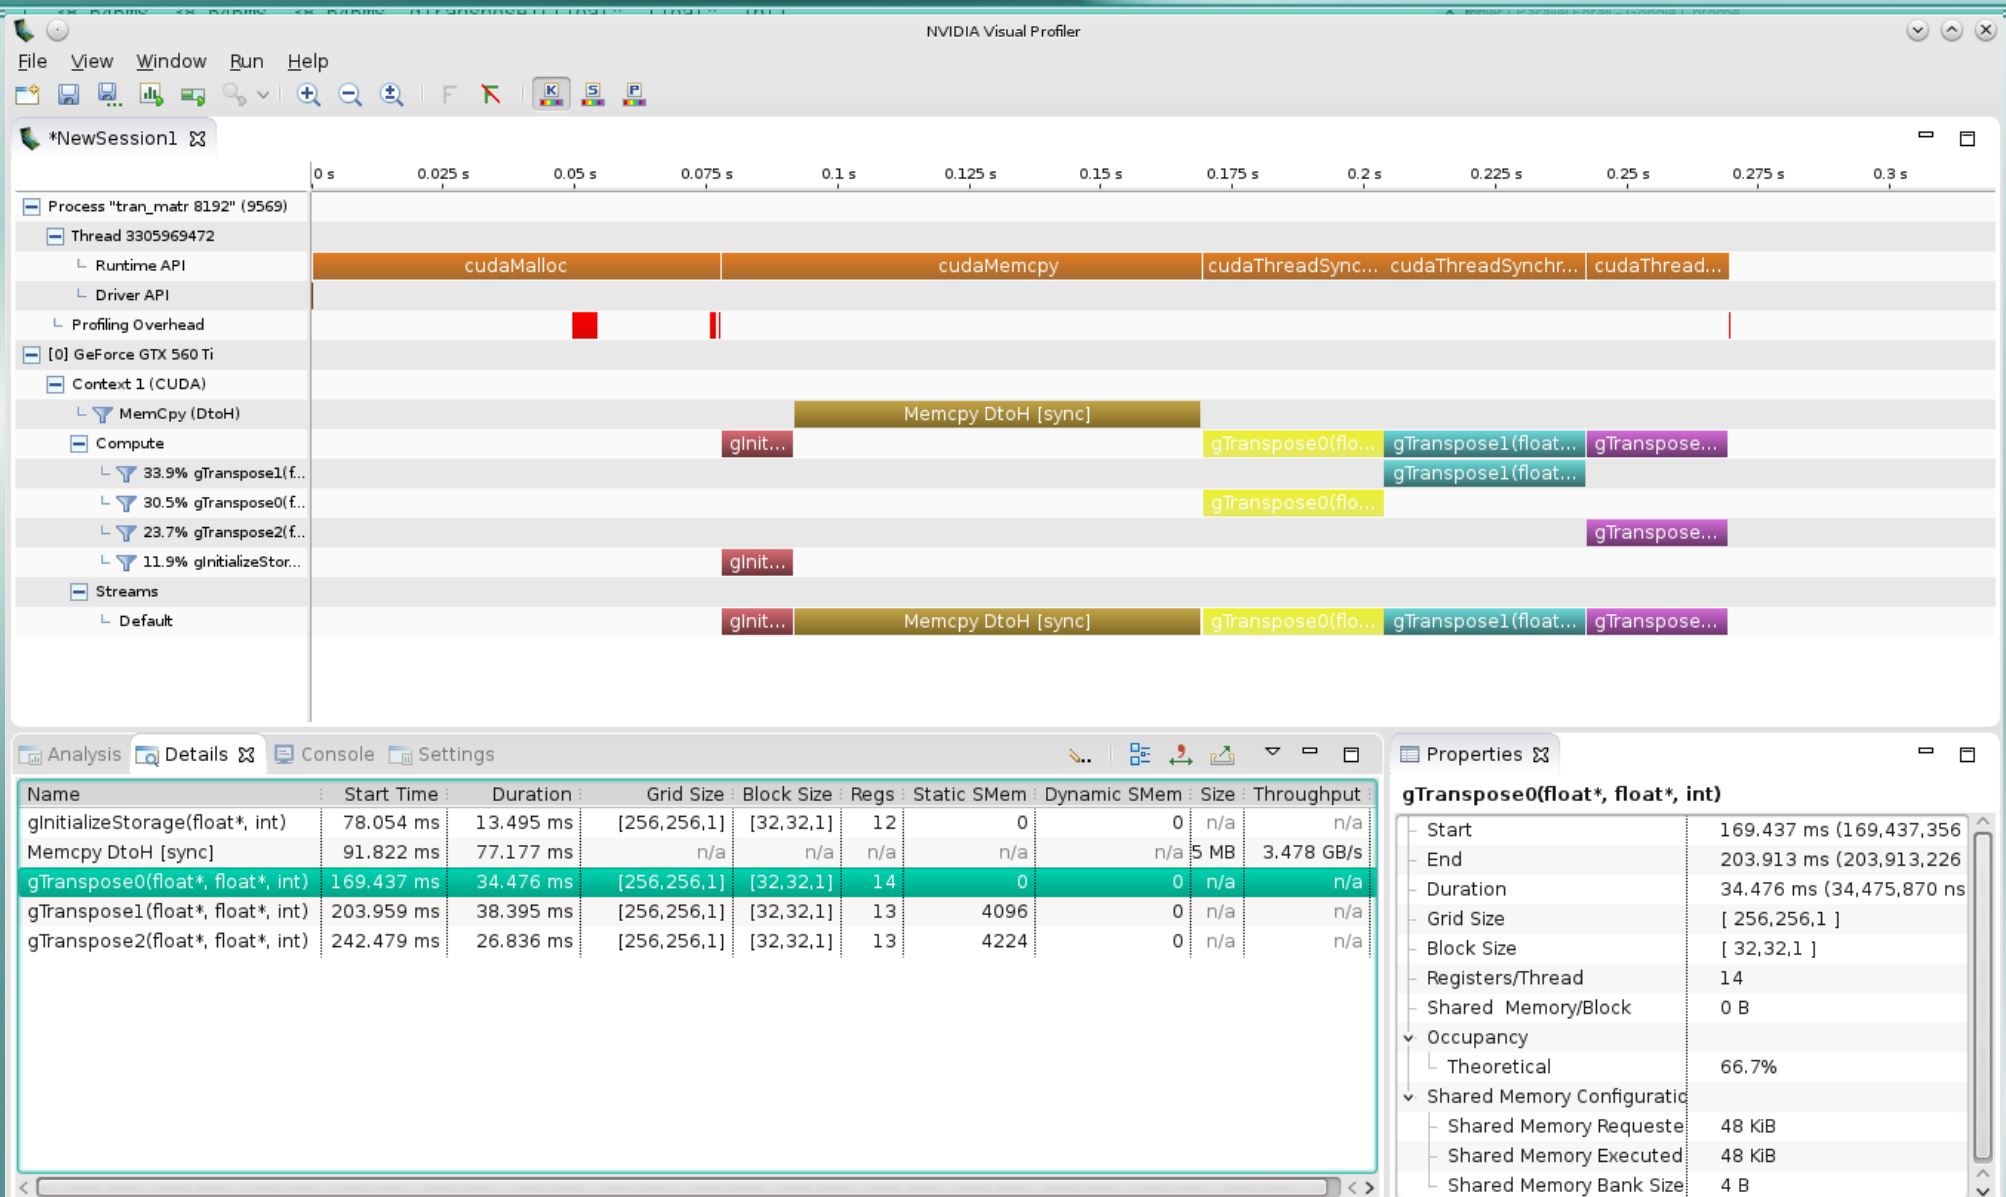
## Время выполнения ядра с разными архитектурами CUDA (Tesla K40m, *Kepler*, CUDA 7.5, *compute capabilities 3.5*)

```
malkov@master:~/WORKSPACE/cuda-gdb/ nvcc -arch=sm_35 tran_matr.cu -o tran_matr
malkov@master:~/WORKSPACE/cuda-gdb/ qsub -I
```

```
.....
      Avg      Min      Max  Name
73.764ms  26.145ms  121.38ms [CUDA memcpy DtoH]
8.3598ms  8.3598ms  8.3598ms gTranspose0(float*, float*, int)
7.5069ms  7.5069ms  7.5069ms gTranspose1(float*, float*, int)
5.9834ms  5.9834ms  5.9834ms gTranspose2(float*, float*, int)
2.5776ms  2.5776ms  2.5776ms gInitializeStorage(float*, int)
```



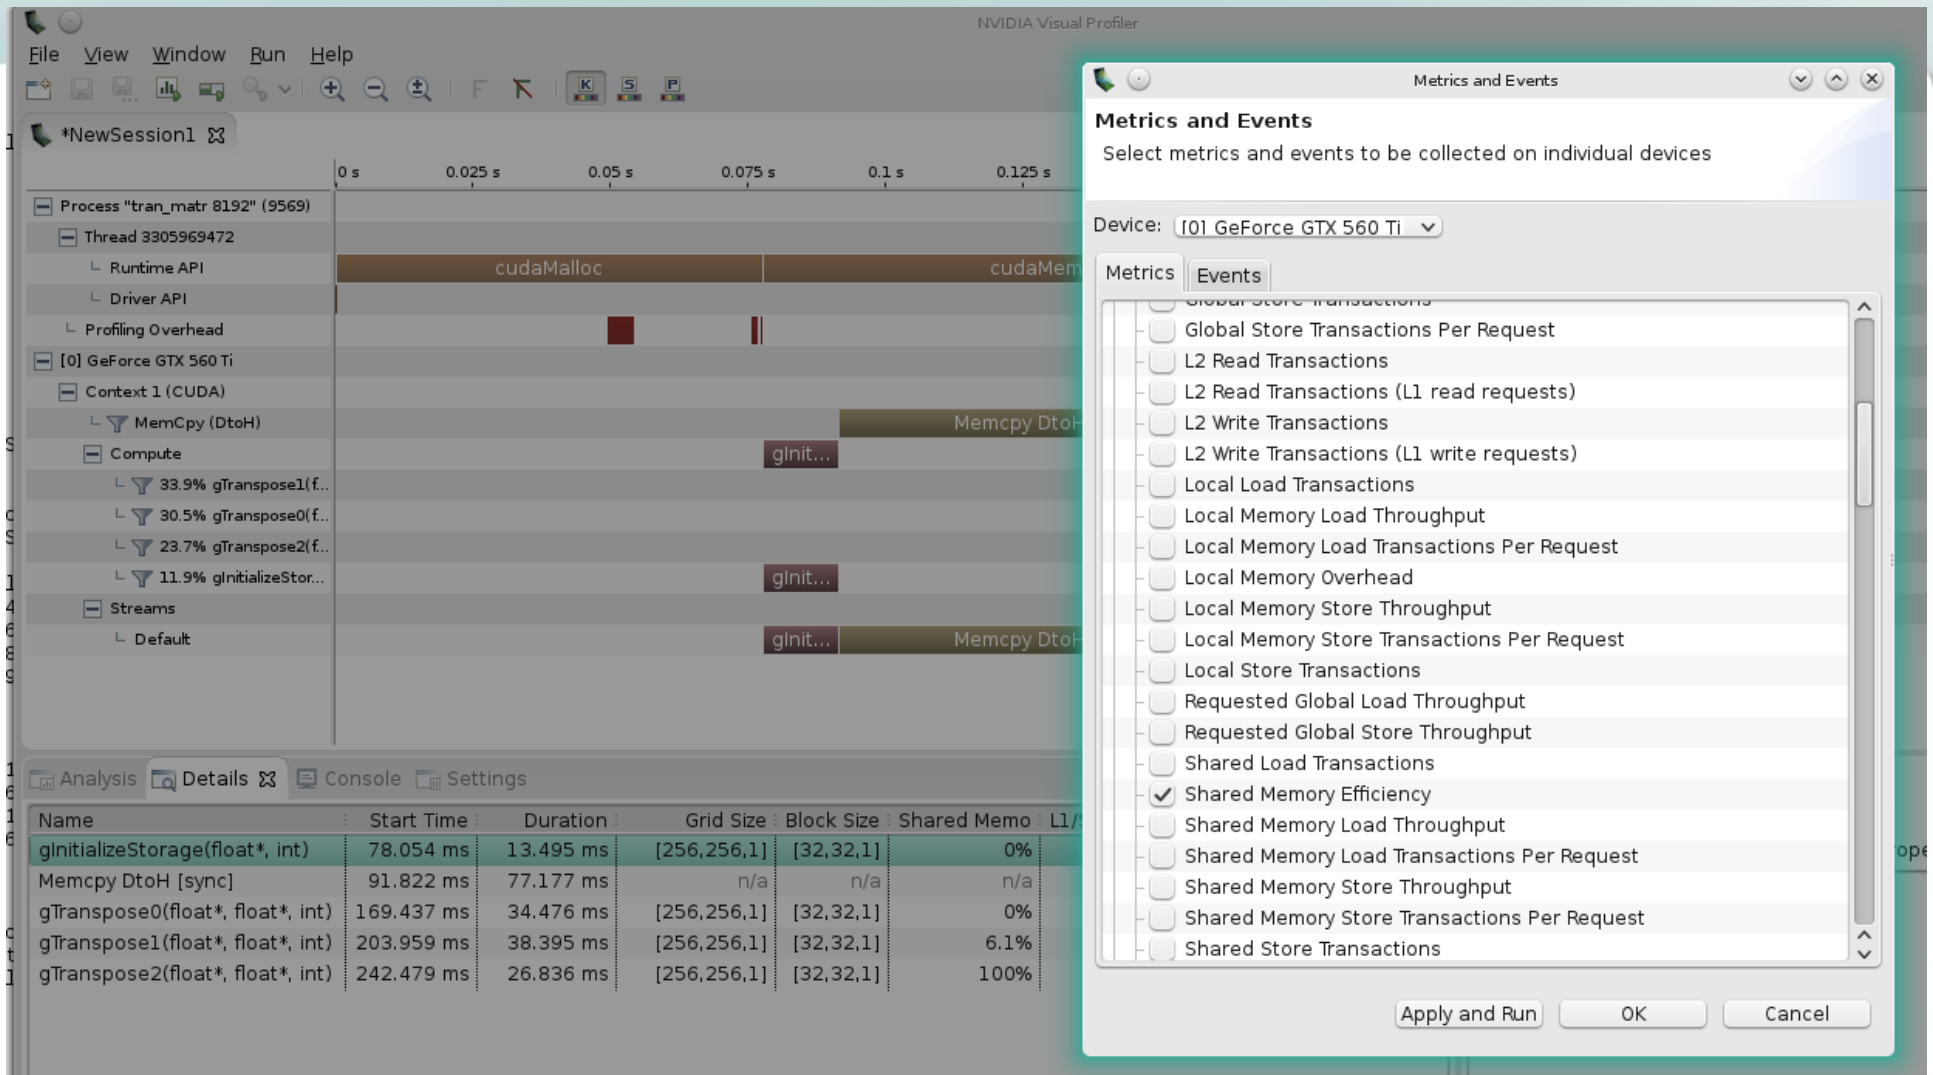
# NVidia Visual Profiler (nvvp)



## *NVidia Visual Profiler (nvvp)*

Name	Start Time	Duration	Grid Size	Block Size
gInitializeStorage(float*, int)	78.054 ms	13.495 ms	[256,256,1]	[32,32,1]
Memcpy DtoH [sync]	91.822 ms	77.177 ms	n/a	n/a
gTranspose0(float*, float*, int)	169.437 ms	34.476 ms	[256,256,1]	[32,32,1]
gTranspose1(float*, float*, int)	203.959 ms	38.395 ms	[256,256,1]	[32,32,1]
gTranspose2(float*, float*, int)	242.479 ms	26.836 ms	[256,256,1]	[32,32,1]

# Добавление событий и счетчиков (nvvp)



**Metrics and Events**

Select metrics and events to be collected on individual devices

Device: [0] GeForce GTX 560 Ti

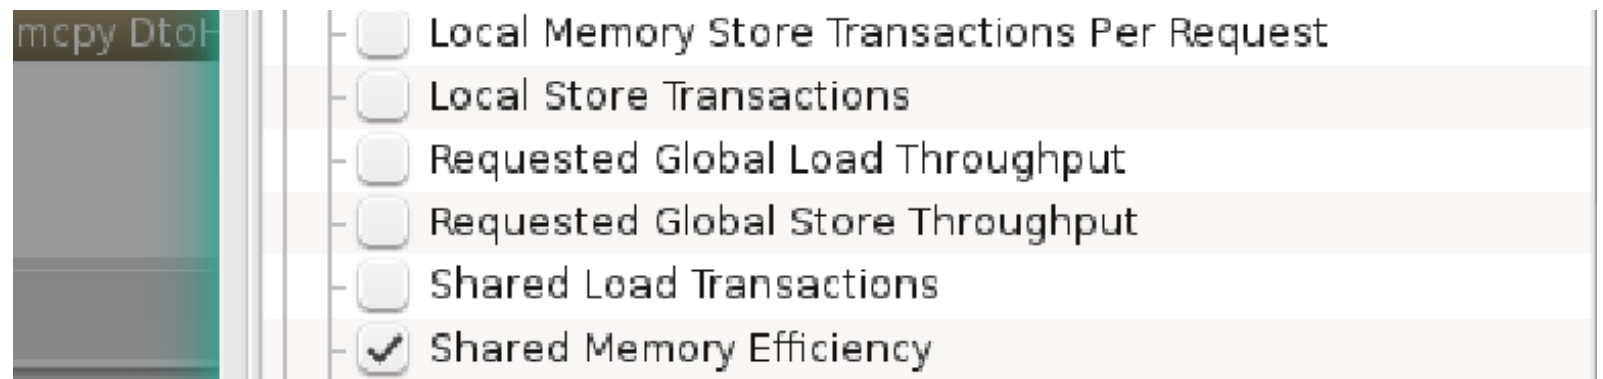
Metrics Events

- ☐ Global Store Transactions
- ☐ Global Store Transactions Per Request
- ☐ L2 Read Transactions
- ☐ L2 Read Transactions (L1 read requests)
- ☐ L2 Write Transactions
- ☐ L2 Write Transactions (L1 write requests)
- ☐ Local Load Transactions
- ☐ Local Memory Load Throughput
- ☐ Local Memory Load Transactions Per Request
- ☐ Local Memory Overhead
- ☐ Local Memory Store Throughput
- ☐ Local Memory Store Transactions Per Request
- ☐ Local Store Transactions
- ☐ Requested Global Load Throughput
- ☐ Requested Global Store Throughput
- ☐ Shared Load Transactions
- ☒ Shared Memory Efficiency
- ☐ Shared Memory Load Throughput
- ☐ Shared Memory Load Transactions Per Request
- ☐ Shared Memory Store Throughput
- ☐ Shared Memory Store Transactions Per Request
- ☐ Shared Store Transactions

Apply and Run OK Cancel

Name	Start Time	Duration	Grid Size	Block Size	Shared Memo	L1/2
glnitializeStorage(float*, int)	78.054 ms	13.495 ms	[256,256,1]	[32,32,1]	0%	
Memcpy DtoH [sync]	91.822 ms	77.177 ms	n/a	n/a	n/a	
gTranspose0(float*, float*, int)	169.437 ms	34.476 ms	[256,256,1]	[32,32,1]	0%	
gTranspose1(float*, float*, int)	203.959 ms	38.395 ms	[256,256,1]	[32,32,1]	6.1%	
gTranspose2(float*, float*, int)	242.479 ms	26.836 ms	[256,256,1]	[32,32,1]	100%	

## Добавление событий и счетчиков (nvvp)



Duration	Grid Size	Block Size	Shared Memo
13.495 ms	[256,256,1]	[32,32,1]	0%
77.177 ms	n/a	n/a	n/a
34.476 ms	[256,256,1]	[32,32,1]	0%
38.395 ms	[256,256,1]	[32,32,1]	6.1%
26.836 ms	[256,256,1]	[32,32,1]	100%

# Добавление событий и счетчиков (nvprof)

```
malkov@linux-5002:~/WORKSHOP/EDUCATION/SibsUTIS/COURSES/2016-2017/CUDA/Lectures/Lecture4/Lab4-2> nvprof --query-metrics | grep shared
shared_load_transactions: Number of shared memory load transactions
shared_store_transactions: Number of shared memory store transactions
shared_load_transactions_per_request: Average number of shared memory load transactions performed for each shared memory load
shared_store_transactions_per_request: Average number of shared memory store transactions performed for each shared memory store
shared_load_throughput: Shared memory load throughput
shared_store_throughput: Shared memory store throughput
shared_efficiency: Ratio of requested shared memory throughput to required shared memory throughput
shared_replay_overhead: Average number of replays due to shared memory conflicts for each instruction executed
ldst_issued: Number of issued local, global, shared and texture memory load and store instructions
ldst_executed: Number of executed local, global, shared and texture memory load and store instructions
l1_shared_utilization: The utilization level of the L1/shared memory relative to peak utilization
ldst_fu_utilization: The utilization level of the multiprocessor function units that execute global, local and shared memory instructions
malkov@linux-5002:~/WORKSHOP/EDUCATION/SibsUTIS/COURSES/2016-2017/CUDA/Lectures/Lecture4/Lab4-2> nvprof --metrics "shared_efficiency" --profile-from-start off ./tran_matr 8192
==10311== NVPROF is profiling process 10311, command: ./tran_matr 8192
    0          2048          4096          6144
16777216      16779264      16781312      16783360
33554432      33556480      33558528      33560576
50331648      50333696      50335744      50337792

==10311== Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.
==10311== Replaying kernel "gTranspose0(float*, float*, int)" (done)
==10311== Replaying kernel "gTranspose1(float*, float*, int)" (done)
==10311== Replaying kernel "gTranspose2(float*, float*, int)" (done)
    0          16777216      33554432      50331648
    2048      16779264      33556480      50333696
    4096      16781312      33558528      50335744
    6144      16783360      33560576      50337792

==10311== Profiling application: ./tran_matr 8192
==10311== Profiling result:
==10311== Metric result:
Invocations      Metric Name      Metric Description      Min      Max      Avg
Device "GeForce GTX 560 Ti (0)"
Kernel: gTranspose2(float*, float*, int)
    1      shared_efficiency      Shared Memory Efficiency      100.00%      100.00%      100.00%
Kernel: gTranspose0(float*, float*, int)
    1      shared_efficiency      Shared Memory Efficiency      0.00%      0.00%      0.00%
Kernel: gTranspose1(float*, float*, int)
    1      shared_efficiency      Shared Memory Efficiency      6.06%      6.06%      6.06%
```

## Добавление событий и счетчиков (nvprof)

```
> nvprof --metrics "shared_efficiency" ./tran_matr 8192
```

Metric Description	Min
Shared Memory Efficiency	100.00%
Shared Memory Efficiency	0.00%
Shared Memory Efficiency	6.06%